

Image Processing

lab 3

Klaas Kliffen Jan Kramer

December 10, 2015

Exercise 1 – 1-D wavelet transforms

- a. Assuming the length of the input vector is a power of two, it is possible to iterate over the input vector applying the algorithm described by example 7.8. The current length is initialized to the length of the vector. For each step, this length is halved. Since the input vector of a step is split in half, a part with sums and a part with differences, the next step will only use the sums, as specified by the algorithm. First two vector are created by taking the odd and even element of the input vector until the current length of the scale. Two result vectors are calculated for the sums and the differences. The difference vector is positioned after the sum vector and written to the input vector. The values of the sum and difference vector needs normalization, which is a factor of the square root of 2.

```
1  % function to perform 1D Haar wavelet transform
2  function retval = IPdwt(x,s)
3  sqrt2 = sqrt(2);
4  out = x;
5
6  initl = length(out);
7
8  for i = 1 : s
9      % Get the odd and even elements
10     odds = out(1:2:initl);
11     evens = out(2:2:initl);
12     % Calculate the means and details
13     sums = (odds + evens);
14     diffs = (odds - evens);
15     % Put the new values
16     out(1:initl) = [sums, diffs] / sqrt2;
17     initl /= 2;
18 end
19
20 retval = out;
21 end
```

- b. For the inverse wavelet transform, the initial length is set to the end of the length of the last sum vector. Then for each step the sum and difference vectors are retrieved from the input vector. Two new component vectors are created for the values. These need to be scaled again to retrieve the right results. The component vectors are then interleaved and

replace their original values in the input vector. This is repeated until the original scale of the transformation is reached.

```
1  % function to perform 1D Haar wavelet transform
2  function retval = IPidwt(x,s)
3  sqrt2 = sqrt(2);
4  out = x;
5
6  % Determine the initial length
7  initl = length(x) / (2^(s-1));
8
9  for i = 1 : s
10     % Retrieve the sums and the differences
11     sums = out(1:initl/2);
12     diffs = out(initl/2+1:initl);
13     % Calculate and scale the result
14     plus = (sums+diffs)/sqrt2;
15     mins = (sums-diffs)/sqrt2;
16     % Combine the new values
17     combined = zeros(initl,1);
18     combined(1:2:end) = plus;
19     combined(2:2:end) = mins;
20     % Store them in the output matrix
21     out(1:initl) = combined;
22     % Increase the length for the next iteration
23     initl *= 2;
24 end
25
26 retval = out;
27 end
```

Exercise 2 – 2-D wavelet transforms

a. TODO: text

```
1  % function to perform 2D Haar wavelet transform
2  function retval = IPdwt2(x, j)
3  % note that x should be double instead of uint, because
4  % the result can get negative
5  out = x;
6  coef = 1/2;
7
8  initrow = size(out, 1);
9  initcol = size(out, 2);
10
11  for i = 1 : j
12      odds_c = out(1:initrow, 1:2:initcol);
13      evens_c = out(1:initrow, 2:2:initcol);
14      sums = (odds_c + evens_c);
15      diffs = (odds_c - evens_c);
16      out(1:initrow, 1:initcol) = [sums, diffs] * coef;
17
18      odds_r = out(1:2:initrow, 1:initcol);
19      evens_r = out(2:2:initrow, 1:initcol);
20      sums = (odds_r + evens_r) * coef;
21      diffs = (odds_r - evens_r) * coef;
22
23      %
24      mid_r = initrow / 2;
25      mid_c = initcol / 2;
26      % approximation image
27      out(1:mid_r, 1:mid_c) = sums(:, 1:mid_c);
28      % vertical detail
29      out(mid_r+1:initrow, 1:mid_c) = sums(:, mid_c+1:initcol);
30      % horizontal detail
31      out(1:mid_r, mid_c+1:initcol) = diffs(:, 1:mid_c);
32      % detail detail
33      out(mid_r+1:initrow, mid_c+1:initcol) = diffs(:, mid_c+1:initcol);
34
35      initrow = mid_r;
36      initcol = mid_c;
37  end
38
39  retval = out;
40
41  endfunction
```

b. TODO: text

```
1  function retval = IPdwt2scale(x, j)
2  % calculate the dwt with a shifted image around 0 (assumes doubles)
3  out = x - 0.5;
4  out = IPdwt2(out, j);
5  out = out + 0.5;
6
```

```

7  initrow = size(out, 1);
8  initcol = size(out, 2);
9
10 % iterate through the levels in the image
11 for i = 1 : j
12     mid_r = initrow / 2;
13     mid_c = initcol / 2;
14
15     % contrast stretch the horizontal details
16     w = out(1:mid_r, mid_c+1:initcol);
17     out(1:mid_r, mid_c+1:initcol) = (w - min(min(w))) * (1 / (max(max(w)
18         )) - min(min(w))));
19
20     % contrast stretch the vertical details
21     w = out(mid_r+1:initrow, 1:mid_c);
22     out(mid_r+1:initrow, 1:mid_c) = (w - min(min(w))) * (1 / (max(max(w)
23         )) - min(min(w))));
24
25     % contrast stretch the diagonal details
26     w = out(mid_r+1:initrow, mid_c+1:initcol);
27     out(mid_r+1:initrow, mid_c+1:initcol) = (w - min(min(w))) * (1 / (
28         max(max(w)) - min(min(w))));
29
30     initrow = mid_r;
31     initcol = mid_c;
32 end
33
34 % contrast stretch the approximation image
35 w = out(1:initrow, 1:initcol);
36 out(1:initrow, 1:initcol) = (w - min(min(w))) * (1 / (max(max(w)) -
37     min(min(w))));
38
39 retval = out;
40 endfunction

```

c. TODO: text

```

1
2 x = im2double(imread(' ../images/vase.tif'));
3
4 %% lab 2 ex 2abc
5 y = IPdwt2(x, 3);
6 imwrite(y, 'unscaled.png');
7 imwrite(im2uint8(IPdwt2scale(x, 3)), 'scaled.png');
8 imwrite(im2uint8(IPdwt2(y,3)), 'output.png');

```




Figure 1: 3 scale wavelet transform of the original image

d. TODO: text




Figure 2: 3 scale wavelet transform restoration

Exercise 3 – Image Compression

- a. First the wavelet transform is performed on the input image. A threshold matrix consisting solely of the threshold value is used to find all pixels larger than the threshold in the transform. The thresholded image is then retrieved by pointwise multiplication of the results of comparing the threshold matrix with the wavelet transform. Since the approximation of the original matrix is also passed by the threshold, it needs to be reconstructed. This is done by copying the values from the wavelet transform back to the thresholded image. The compressed image is then retrieved by applying the inverse wavelet transform.

The root mean square error and the mean square signal to noise ratio are calculated from the given formulas. For the compression ratio the the build-in function of entropy is used to calculate the entropy of the compressed and original image. The entropy is a value for how complex an image is. Compressing the image lowers the complexity. Dividing the original entropy by the compressed entropy yields a compression ratio.

```

1  % Function to compress an image using wavelet transform
2  function retval = IPwaveletcompress(img, scale, threshold)
3
4  [width,height] = size(img);
5  wl = width / (2^scale);
6  hl = height / (2^scale);
7
8  % Wavelet transform
9  wtrans = IPdwt2(img,scale);
10
11 % Construct a matrix for the threshold
12 thresholdmat = threshold * ones(size(img));
13 % Matrix containing 1 for pixels above the threshold
14 results = abs(wtrans) > thresholdmat;
15 % Perform the thresholding by elementwise multiplying
16 threshed = zeros(size(img));
17 threshed = wtrans .*results;
18
19 % Copy the original image part (for dark values in the original
20 threshed(1:hl,1:wl) = wtrans(1:hl,1:wl);
21
22 % Convert it back
23 compressed = IPidwt2(threshed, scale);
24
25 printf("Scale: %d Threshold: %f\n", scale, threshold);
26 error = compressed - img;
27 errorsq = error .* error;
28 rmse = sqrt(mean(mean(errorsq)));
29 printf("Root mean square error: %f\n",rmse);
30
31 squared = compressed .* compressed;
32 snr = sum(sum(squared)) / sum(sum(errorsq));
33 printf("Mean square signal to noise: %f\n",snr);
34
35 % Calculatute the compression
36 orig = entropy(im2uint8(img));
37 comp = entropy(im2uint8(threshed));
38 printf("Compress ratio: %f:1\n\n",orig/comp);
39
40
41 retval = compressed;
42
43 end


```

- b. The quantitative properties for several different scales and thresholds can be seen in table 1. Globally the signal to noise ratio decreases and the errors and compression ratio increase while increasing scale and threshold. Although for this image the increasing the scale


past 9 was not possible, due to its size being a square of 512 pixels. It would seem that increasing it further would not compress the image any further past a compression ratio of 37. Increasing the threshold will yield in higher compression ratios. Although the signal to noise ratio decreases exponentiall, while the errors increase almost linearly.

| Scale | Threshold | ϵ_{rms} | SNR | Compression ratio |
|-------|-----------|------------------|---------|-------------------|
| 1 | 0.02 | 0.0086 | 1277.34 | 2.59:1 |
| 3 | 0.02 | 0.0157 | 385.47 | 17.42:1 |
| 3 | 0.05 | 0.0249 | 152.09 | 24.87:1 |
| 3 | 0.10 | 0.0322 | 90.33 | 29.14:1 |
| 5 | 0.02 | 0.0220 | 194.71 | 34.00:1 |
| 7 | 0.02 | 0.0272 | 127.44 | 36.40:1 |
| 7 | 0.05 | 0.0575 | 27.627 | 113.18:1 |
| 7 | 0.10 | 0.1035 | 7.837 | 460.20:1 |
| 9 | 0.02 | 0.0321 | 91.02 | 36.55:1 |

Table 1: Quantitative compression quality for different scales and threshold



Original image



Scale: 1 Threshold: 0.02

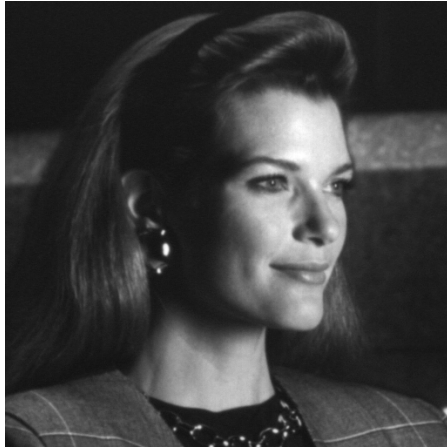


Scale: 5 Threshold: 0.02



Scale: 9 Threshold: 0.02

Figure 3: Increasing wavelet compression scale on an image



Original image



Scale: 3 Threshold: 0.02




Scale: 3 Threshold: 0.05



Scale: 3 Threshold: 0.10

Figure 4: Increasing wavelet compression threshold on scale 3 on an image




Original image



Scale: 7 Threshold: 0.02



Scale: 7 Threshold: 0.05



Scale: 7 Threshold: 0.10

Figure 5: Increasing wavelet compression threshold on scale 7 on an image

Task distribution

| ex1 | design | implementation | answers questions | writing report |
|-------|--------|----------------|-------------------|----------------|
| Klaas | 60% | 90% | n.a. | 75% |
| Jan | 40% | 10% | n.a. | 25% |

| ex2 | design | implementation | answers questions | writing report |
|-------|--------|----------------|-------------------|----------------|
| Klaas | 40% | 20% | 50% | 25% |
| Jan | 60% | 80% | 50% | 75% |

| ex3 | design | implementation | answers questions | writing report |
|-------|--------|----------------|-------------------|----------------|
| Klaas | 50% | 75% | 50% | 75% |
| Jan | 50% | 25% | 50% | 25% |